

```
In [1]: # 1. IMPORT REQUIRED LIBRARIES
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
```

```
In [2]: # 2. LOAD THE TRAINING DATASET
```

	Unnamed:	SeriousDlqn2yrs	RevolvingUtilizationOfUnsecuredLines	age	NumberofTime30-59DaysPastDueNotWorse	DebtRatio	MonthlyIncome	NumberofOpenCreditLinesAndLoans	NumberofTimes90DaysLate	NumberRealEstateLoansOrLines	NumberofTime60-89DaysPastDueNotWorse	NumberofDependents
0	1	1	0.766127	45	2	0.802982	9120.0	13	0	6	0	2.0
1	2	0	0.957151	40	0	0.121876	2800.0	4	0	0	0	1.0
2	3	0	0.658180	38	1	0.085113	3042.0	2	1	0	0	0.0
3	4	0	0.233810	30	0	0.036050	3300.0	5	0	0	0	0.0
4	5	0	0.907239	49	1	0.024926	63588.0	7	0	1	0	0.0

```
In [3]: # 3. BASIC DATA CLEANING
```

```
# Drop unnamed index column if exists
df = df.loc[:, ~df.columns.str.contains('Unnamed')]

# Check missing values
df.isnull().sum()
```

```
Out[3]:
```

```
SeriousDlqn2yrs      0
RevolvingUtilizationOfUnsecuredLines  0
age                  0
NumberofTime30-59DaysPastDueNotWorse  0
DebtRatio             0
MonthlyIncome         29731
NumberOfOpenCreditLinesAndLoans        0
NumberofTimes90DaysLate               0
NumberRealEstateLoansOrLines          0
NumberofTimes60-89DaysPastDueNotWorse  0
NumberofDependents                 3924
dtype: int64
```

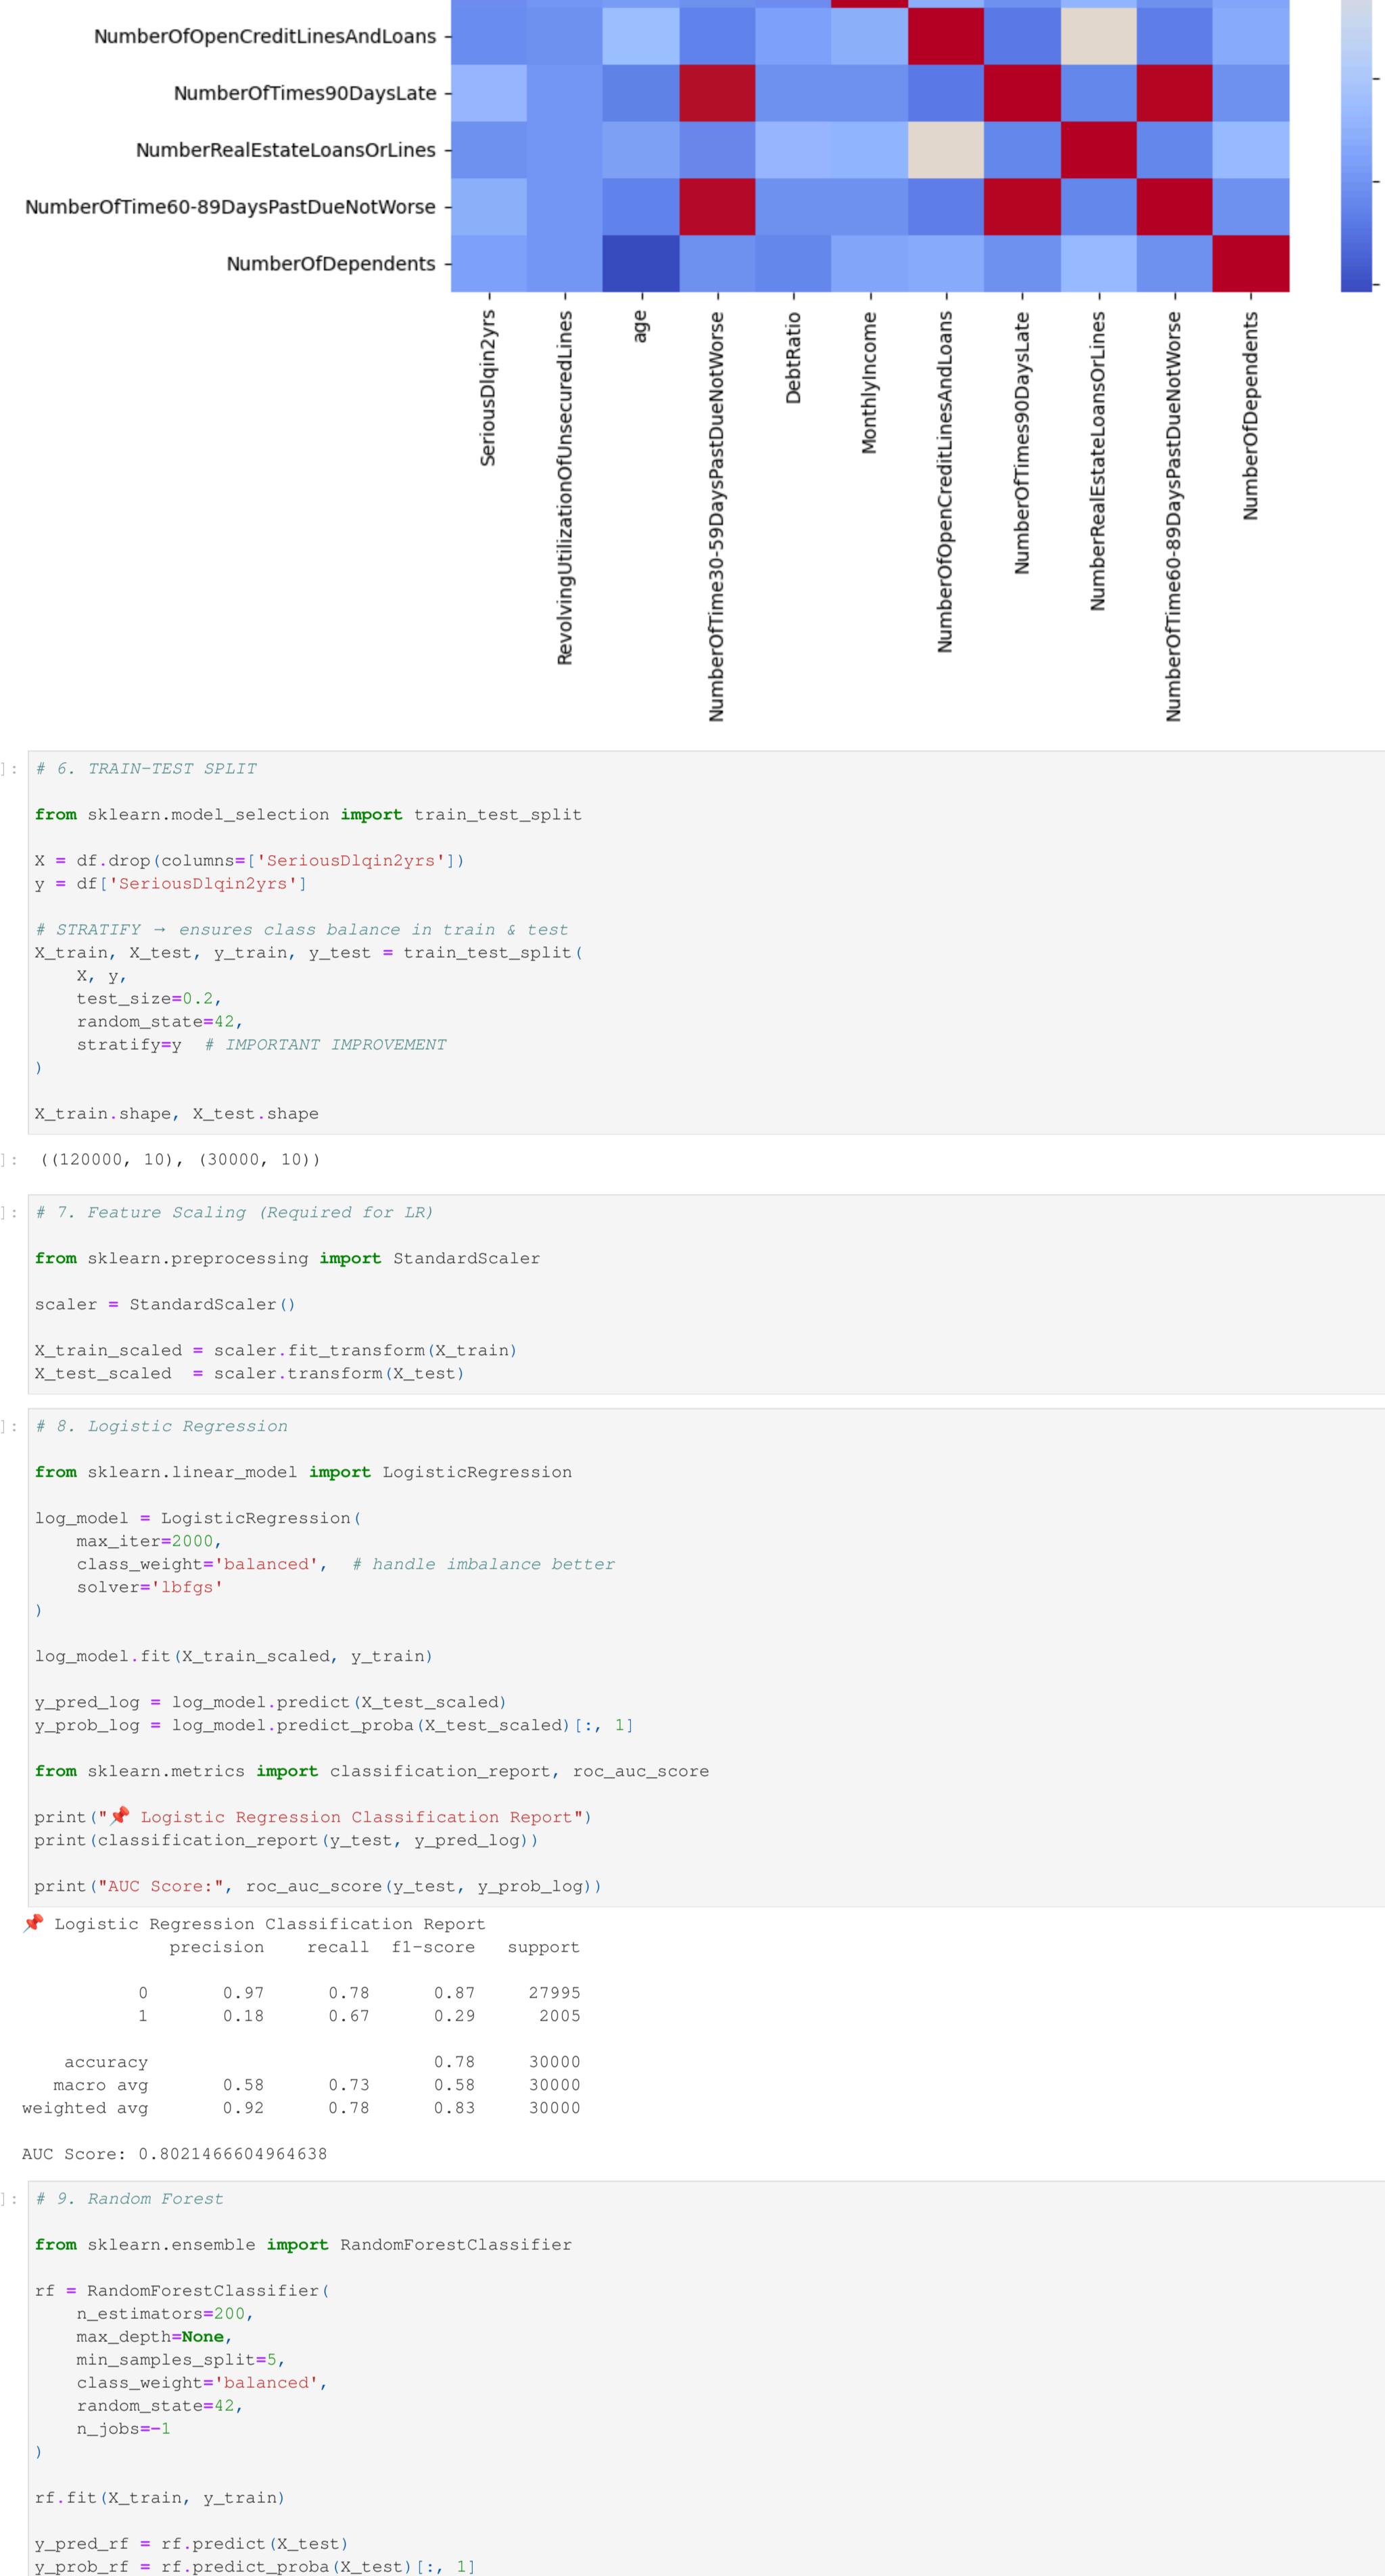
```
In [4]: # 4. MonthlyIncome & NumberofDependents have missing values (also fill with median)
```

```
df['MonthlyIncome'] = df['MonthlyIncome'].fillna(df['MonthlyIncome'].median())
df['NumberofDependents'] = df['NumberofDependents'].fillna(df['NumberofDependents'].median())
```

```
In [5]: # 5. Target distribution
```

```
sns.countplot(x=df['SeriousDlqn2yrs'])
plt.title("Target Distribution")
plt.show()

# Correlation heatmap
plt.figure(figsize=(10,4))
sns.heatmap(df.corr(), cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```



```
In [6]: # 6. TRAIN-TEST SPLIT
```

```
from sklearn.model_selection import train_test_split
x = df.drop(columns=['SeriousDlqn2yrs'])
y = df['SeriousDlqn2yrs']

# STRATIFY - ensures class balance in train & test
X_train, X_test, y_train, y_test = train_test_split(
    x, y, test_size=0.2,
    random_state=42,
    stratify=y # IMPORTANT IMPROVEMENT
)
```

```
X_train.shape, X_test.shape
```

```
Out[6]: ((120000, 10), (30000, 10))
```

```
In [7]: # 7. Feature Scaling (Required for LR)
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [8]: # 8. Logistic Regression
```

```
from sklearn.linear_model import LogisticRegression
```

```
log_model = LogisticRegression(
    max_iter=2000,
    class_weight='balanced', # handle imbalance better
    solver='lbfgs'
)
```

```
log_model.fit(X_train_scaled, y_train)
```

```
y_pred_log = log_model.predict(X_test_scaled)
y_prob_log = log_model.predict_proba(X_test_scaled)[:, 1]
```

```
from sklearn.metrics import classification_report, roc_auc_score
```

```
print("Logistic Regression Classification Report")
```

```
print(classification_report(y_test, y_pred_log))
```

```
print("AUC Score:", roc_auc_score(y_test, y_prob_log))
```

```
* Logistic Regression Classification Report
```

```
precision    recall   f1-score   support
```

```
0       0.97      0.78      0.87     27995
```

```
1       0.18      0.67      0.29     2005
```

```
accuracy           0.78     30000
```

```
macro avg       0.58      0.73      0.58     30000
```

```
weighted avg     0.92      0.78      0.89     30000
```

```
AUC Score: 0.8221466604964638
```

```
In [9]: # 9. Random Forest
```

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(
    n_estimators=200,
    max_depth=None,
    min_samples_split=5,
    class_weight='balanced',
    random_state=42,
    n_jobs=-1
)
```

```
rf.fit(X_train, y_train)
```

```
y_pred_rf = rf.predict(X_test)
y_prob_rf = rf.predict_proba(X_test)[:, 1]
```

```
print("Random Forest Classification Report")
```

```
print(classification_report(y_test, y_pred_rf))
```

```
print("AUC Score:", roc_auc_score(y_test, y_prob_rf))
```

```
* Random Forest Classification Report
```

```
precision    recall   f1-score   support
```

```
0       0.95      0.98      0.97     27995
```

```
1       0.51      0.24      0.32     2005
```

```
accuracy           0.93     30000
```

```
macro avg       0.73      0.61      0.65     30000
```

```
weighted avg     0.92      0.93      0.92     30000
```

```
AUC Score: 0.847598321479631
```

```
In [10]: # 10. Confusion Matrix
```

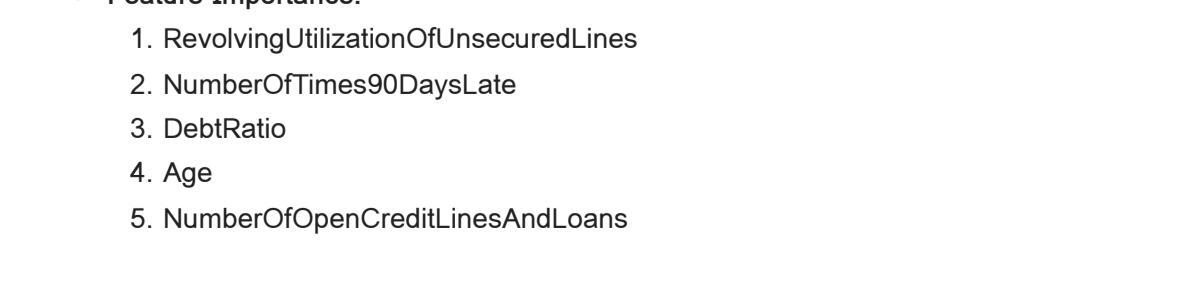
```
from sklearn.metrics import confusion_matrix
import seaborn as sns
```

```
cm = confusion_matrix(y_test, y_pred_rf)
cm_norm = cm / cm.sum()
```

```
plt.figure(figsize=(6,5))
sns.heatmap(cm_norm, annot=True, fmt=".2%", cmap="Blues")
plt.title("Normalized Confusion Matrix - Random Forest")
plt.xlabel("Predicted")
plt.ylabel("Actual")
```

```
plt.show()
```

Normalized Confusion Matrix - Random Forest



```
In [11]: # 11. Feature Importance (Top 10)
```

```
import pandas as pd
import matplotlib.pyplot as plt
```

```
importances = rf.feature_importances_
features = X.columns
```

```
fi_df = pd.DataFrame({
    "Feature": features,
    "Importance": importances
}).sort_values("Importance", ascending=False).head(10)
```

```
plt.figure(figsize=(8,5))
plt.barh(fi_df['Feature'], fi_df['Importance'])
plt.gca().invert_yaxis()
plt.title("Top 10 Important Features (Random Forest)")
plt.xlabel("Importance")
plt.show()
```

Top 10 Important Features (Random Forest)



```
In [12]: # 12. Final Model Selection
```

```
rf_auc = roc_auc_score(y_test, rf.predict_proba(X_test)[:, 1])
log_auc = roc_auc_score(y_test, log_model.predict_proba(X_test_scaled)[:, 1])
```

```
print("Random Forest AUC:", rf_auc)
print("Logistic Regression AUC:", log_auc)
```

```
if rf_auc > log_auc:
    print("\n\nFinal Model Selected: Random Forest")
else:
    print("\n\nFinal Model Selected: Logistic Regression")
```

```
Random Forest AUC: 0.847598321479631
Logistic Regression AUC: 0.8221466604964638
```

Final Model Selected: Random Forest

```
In [13]: # 13. Export Final Model
```

```
import joblib, sys
```

```
joblib.dump(rf, "credit_risk_model.pkl")
print("Model Saved Successfully!")
```

```
Model Saved Successfully!
```

```
In [14]: # 14. SHAP Explainability (Highly Recommended)
```

```
import shap
import matplotlib.pyplot as plt
```

```
sample = X_test.sample(100, random_state=42)
```

```
explainer = shap.TreeExplainer(rf)
shap_values = explainer.shap_values(sample)
```

```
shap.summary_plot(shap_values, sample, plot_type="bar")
```

SHAP summary plot

SHAP interaction plot

SHAP dependence plot

SHAP partial dependence plot

SHAP coefficient plot

SHAP feature importance plot

SHAP feature contribution plot

SHAP feature distribution plot

SHAP feature summary plot

SHAP feature value plot

SHAP feature weight plot

SHAP model summary plot

SHAP model value plot

SHAP plot

SHAP plot by class

SHAP plot by feature

SHAP plot by label

SHAP plot by prediction

SHAP plot by row

SHAP plot by value

SHAP plot by weight

SHAP plot by feature

SHAP plot by label

SHAP plot by prediction

SHAP plot by row

SHAP plot by value

SHAP plot by weight

SHAP plot by feature

SHAP plot by label

SHAP plot by prediction

SHAP plot by row

SHAP plot by value

SHAP plot by weight

SHAP plot by feature

SHAP plot by label

SHAP plot by prediction

SHAP plot by row

SHAP plot by value

SHAP plot by weight

SHAP plot by feature

SHAP plot by label

SHAP plot by prediction

SHAP plot by row

SHAP plot by value

SHAP plot by weight

SHAP plot by feature

SHAP plot by label

